

```
In [0]: import matplotlib.pyplot as plt
from keras.utils import np_utils
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization
from keras.optimizers import Adam
from keras.layers import Dense, Activation
import seaborn as sns
import numpy as np
import keras
```

```
%matplotlib inline
```

```
In [2]: (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 0s 0us/step

```
In [3]: print("Training data shape: ", X_train.shape) # (60000, 28, 28) -- 60000 image
s, each 28x28 pixels
print("Test data shape", X_test.shape)

print("Training label shape: ", y_train.shape) # (60000,) -- 60000 numbers (al
l 0-9)
print("First 5 training labels: ", y_train[:5]) # [5, 0, 4, 1, 9]
```

Training data shape: (60000, 28, 28)  
Test data shape (10000, 28, 28)  
Training label shape: (60000,)  
First 5 training labels: [5 0 4 1 9]

```
In [0]: # Flatten the images
image_vector_size = 28*28
X_train = X_train.reshape(X_train.shape[0], image_vector_size)
X_test = X_test.reshape(X_test.shape[0], image_vector_size)
```

```
In [5]: # Convert to "one-hot" vectors using the to_categorical function
num_classes = 10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
print("First 5 training lables as one-hot encoded vectors:\n", y_train[:5])
```

First 5 training lables as one-hot encoded vectors:  
[[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]  
[1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]  
[0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]  
[0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]  
[0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]]

```
In [0]:
```

## Model 1 without dropout and BN

```
In [0]: output_dim = 10
input_dim = X_train.shape[1]

nb_epoch = 10
batch_size = 128
```

```
In [0]: model1 = Sequential()

model1.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model1.add(Dense(120, activation='relu', kernel_initializer='random_uniform'))

model1.add(Dense(output_dim, activation='softmax'))

print(model1.summary())
```

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 364)	285740
dense_16 (Dense)	(None, 120)	43800
dense_17 (Dense)	(None, 10)	1210
Total params: 330,750		
Trainable params: 330,750		
Non-trainable params: 0		
None		

```
In [0]: model1.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
history = model1.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

Epoch 2/20

Epoch 3/20

Epoch 4/20

Epoch 5/20

Epoch 6/20

Epoch 7/20

Epoch 8/20

Epoch 9/20

Epoch 10/20

Epoch 11/20

Epoch 12/20

Epoch 13/20

Epoch 14/20

Epoch 15/20

Epoch 16/20

Epoch 17/20

Epoch 18/20

Epoch 19/20

Epoch 20/20

```
In [0]: score = model1.evaluate(X_test, y_test)
```


10000/10000 [=====] - 0s 49us/step

```
In [0]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.8604

## 2 Layers with BN and Dropout

```
In [0]: model2 = Sequential()

model2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model2.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))

model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(output_dim, activation='softmax'))

print(model2.summary())
```

Layer (type)	Output Shape	Param #
dense_21 (Dense)	(None, 364)	285740
dense_22 (Dense)	(None, 128)	46720
batch_normalization_3 (Batch Normalization)	(None, 128)	512
dropout_3 (Dropout)	(None, 128)	0
dense_23 (Dense)	(None, 10)	1290
Total params: 334,262		
Trainable params: 334,006		
Non-trainable params: 256		

None

```
In [0]: model2.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

Epoch 2/20

Epoch 3/20

Epoch 4/20

Epoch 5/20

Epoch 6/20

Epoch 7/20

Epoch 8/20

Epoch 9/20

Epoch 10/20

Epoch 11/20

Epoch 12/20

Epoch 13/20

Epoch 14/20

Epoch 15/20

Epoch 16/20

Epoch 17/20

Epoch 18/20

Epoch 19/20

Epoch 20/20

```
In [0]: score = model2.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 55us/step

```
In [0]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9819

## Changing dropout rate to 0.8

```
In [16]: model2 = Sequential()

model2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model2.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))

model2.add(BatchNormalization())
model2.add(Dropout(0.8))

model2.add(Dense(output_dim, activation='softmax'))

print(model2.summary())
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version. Instructions for updating:  
Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 364)	285740
dense_2 (Dense)	(None, 128)	46720
batch_normalization_1 (Batch Normalization)	(None, 128)	512
dropout_1 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290
Total params: 334,262		
Trainable params: 334,006		
Non-trainable params: 256		

None

```
In [0]: model2.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```



```
In [19]: # Training the model  
history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1610 -  
acc: 0.9561 - val\_loss: 0.1074 - val\_acc: 0.9673

Epoch 2/20

60000/60000 [=====] - 6s 92us/step - loss: 0.1238 -  
acc: 0.9661 - val\_loss: 0.0930 - val\_acc: 0.9745

Epoch 3/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1043 -  
acc: 0.9710 - val\_loss: 0.0836 - val\_acc: 0.9748

Epoch 4/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0880 -  
acc: 0.9758 - val\_loss: 0.0896 - val\_acc: 0.9748

Epoch 5/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0741 -  
acc: 0.9792 - val\_loss: 0.0880 - val\_acc: 0.9743

Epoch 6/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0638 -  
acc: 0.9816 - val\_loss: 0.0733 - val\_acc: 0.9808

Epoch 7/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0622 -  
acc: 0.9829 - val\_loss: 0.0842 - val\_acc: 0.9780

Epoch 8/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0528 -  
acc: 0.9850 - val\_loss: 0.0758 - val\_acc: 0.9795

Epoch 9/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0472 -  
acc: 0.9863 - val\_loss: 0.0799 - val\_acc: 0.9796

Epoch 10/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0426 -  
acc: 0.9875 - val\_loss: 0.0946 - val\_acc: 0.9776

Epoch 11/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0409 -  
acc: 0.9878 - val\_loss: 0.0826 - val\_acc: 0.9803

Epoch 12/20

60000/60000 [=====] - 6s 92us/step - loss: 0.0345 -  
acc: 0.9894 - val\_loss: 0.0852 - val\_acc: 0.9798

Epoch 13/20

60000/60000 [=====] - 5s 92us/step - loss: 0.0312 -  
acc: 0.9905 - val\_loss: 0.0869 - val\_acc: 0.9798

Epoch 14/20

60000/60000 [=====] - 5s 91us/step - loss: 0.0307 -  
acc: 0.9911 - val\_loss: 0.0837 - val\_acc: 0.9815

Epoch 15/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0285 -  
acc: 0.9915 - val\_loss: 0.1004 - val\_acc: 0.9785

Epoch 16/20

60000/60000 [=====] - 5s 89us/step - loss: 0.0288 -  
acc: 0.9913 - val\_loss: 0.1110 - val\_acc: 0.9756

Epoch 17/20

60000/60000 [=====] - 5s 90us/step - loss: 0.0265 -  
acc: 0.9922 - val\_loss: 0.0863 - val\_acc: 0.9815

Epoch 18/20

60000/60000 [=====] - 6s 93us/step - loss: 0.0237 -  
acc: 0.9927 - val\_loss: 0.0986 - val\_acc: 0.9788

Epoch 19/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0234 -

```
acc: 0.9930 - val_loss: 0.0881 - val_acc: 0.9803
Epoch 20/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0185 -
acc: 0.9945 - val_loss: 0.0934 - val_acc: 0.9803
```

```
In [20]: score = model2.evaluate(X_test, y_test)
```


```
10000/10000 [=====] - 0s 50us/step
```

```
In [21]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [22]: print('Test accuracy:', score[1])
```

```
Test accuracy: 0.9803
```

## Changing dropout rate to 0.2

```
In [30]: model2 = Sequential()

model2.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model2.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))

model2.add(BatchNormalization())
model2.add(Dropout(0.2))

model2.add(Dense(output_dim, activation='softmax'))

print(model2.summary())
```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 364)	285740
dense_8 (Dense)	(None, 128)	46720
batch_normalization_3 (Batch Normalization)	(None, 128)	512
dropout_3 (Dropout)	(None, 128)	0
dense_9 (Dense)	(None, 10)	1290
Total params: 334,262		
Trainable params: 334,006		
Non-trainable params: 256		

None

```
In [0]: model2.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [32]: # Training the model
history = model2.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 101us/step - loss: 0.2233 - acc: 0.9340 - val\_loss: 0.1149 - val\_acc: 0.9643

Epoch 2/10

60000/60000 [=====] - 5s 89us/step - loss: 0.0950 - acc: 0.9714 - val\_loss: 0.0874 - val\_acc: 0.9731

Epoch 3/10

60000/60000 [=====] - 5s 89us/step - loss: 0.0646 - acc: 0.9803 - val\_loss: 0.0940 - val\_acc: 0.9713

Epoch 4/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0496 - acc: 0.9841 - val\_loss: 0.0780 - val\_acc: 0.9759

Epoch 5/10

60000/60000 [=====] - 5s 89us/step - loss: 0.0382 - acc: 0.9874 - val\_loss: 0.0642 - val\_acc: 0.9815

Epoch 6/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0315 - acc: 0.9902 - val\_loss: 0.0838 - val\_acc: 0.9747

Epoch 7/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0269 - acc: 0.9912 - val\_loss: 0.0745 - val\_acc: 0.9792

Epoch 8/10

60000/60000 [=====] - 6s 92us/step - loss: 0.0237 - acc: 0.9925 - val\_loss: 0.0871 - val\_acc: 0.9754

Epoch 9/10

60000/60000 [=====] - 6s 92us/step - loss: 0.0215 - acc: 0.9926 - val\_loss: 0.0807 - val\_acc: 0.9770

Epoch 10/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0180 - acc: 0.9941 - val\_loss: 0.0750 - val\_acc: 0.9813

```
In [33]: score = model2.evaluate(X_test, y_test)
```


10000/10000 [=====] - 0s 48us/step

```
In [34]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [35]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9813

## Model with 3 hidden layers

```
In [0]: model3 = Sequential()

model3.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model3.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model3.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))

#model2.add(BatchNormalization())
#model2.add(Dropout(0.5))

model3.add(Dense(output_dim, activation='softmax'))

print(model3.summary())
```

Layer (type)	Output Shape	Param #
dense_28 (Dense)	(None, 364)	285740
dense_29 (Dense)	(None, 128)	46720
dense_30 (Dense)	(None, 64)	8256
dense_31 (Dense)	(None, 10)	650
Total params: 341,366		
Trainable params: 341,366		
Non-trainable params: 0		
None		

```
In [0]: model3.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
history = model3.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

Epoch 2/20

Epoch 3/20

Epoch 4/20

Epoch 5/20

Epoch 6/20

Epoch 7/20

Epoch 8/20

Epoch 9/20

Epoch 10/20

Epoch 11/20

Epoch 12/20

Epoch 13/20

Epoch 14/20

Epoch 15/20

Epoch 16/20

Epoch 17/20

Epoch 18/20

Epoch 19/20

Epoch 20/20

```
In [0]: score = model3.evaluate(X_test, y_test)
```

10000/10000 [=====] - 1s 53us/step




```
In [0]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9774

```
In [0]:
```

```
In [0]:
```

## Model with 3 hidden layers along with BN and dropout

```
In [0]: model4 = Sequential()

model4.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model4.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))

model4.add(BatchNormalization())
model4.add(Dropout(0.5))

model4.add(Dense(output_dim, activation='softmax'))

print(model4.summary())
```

Layer (type)	Output Shape	Param #
dense_32 (Dense)	(None, 364)	285740
dense_33 (Dense)	(None, 128)	46720
dense_34 (Dense)	(None, 64)	8256
batch_normalization_4 (Batch Normalization)	(None, 64)	256
dropout_4 (Dropout)	(None, 64)	0
dense_35 (Dense)	(None, 10)	650
Total params: 341,622		
Trainable params: 341,494		
Non-trainable params: 128		

None

```
In [0]: model4.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

Epoch 2/20

Epoch 3/20

Epoch 4/20

Epoch 5/20

Epoch 6/20

Epoch 7/20

Epoch 8/20

Epoch 9/20

Epoch 10/20

Epoch 11/20

Epoch 12/20

Epoch 13/20

Epoch 14/20

Epoch 15/20

Epoch 16/20

Epoch 17/20

Epoch 18/20

Epoch 19/20

Epoch 20/20

```
In [0]: score = model4.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 58us/step

```
In [0]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9837

**changing dropout rate to 0.8**

```
In [36]: model4 = Sequential()

model4.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model4.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))

model4.add(BatchNormalization())
model4.add(Dropout(0.8))

model4.add(Dense(output_dim, activation='softmax'))

print(model4.summary())
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_10 (Dense)	(None, 364)	285740
dense_11 (Dense)	(None, 128)	46720
dense_12 (Dense)	(None, 64)	8256
batch_normalization_4 (Batch Normalization)	(None, 64)	256
dropout_4 (Dropout)	(None, 64)	0
dense_13 (Dense)	(None, 10)	650
=====	=====	=====
Total params: 341,622		
Trainable params: 341,494		
Non-trainable params: 128		
None		

```
In [0]: model4.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```

```
In [38]: # Training the model
history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 104us/step - loss: 0.7021 - acc: 0.7874 - val\_loss: 0.5703 - val\_acc: 0.9000

Epoch 2/10

60000/60000 [=====] - 5s 90us/step - loss: 0.2923 - acc: 0.9194 - val\_loss: 0.1403 - val\_acc: 0.9595

Epoch 3/10

60000/60000 [=====] - 5s 91us/step - loss: 0.1997 - acc: 0.9452 - val\_loss: 0.1222 - val\_acc: 0.9656

Epoch 4/10

60000/60000 [=====] - 5s 91us/step - loss: 0.1612 - acc: 0.9552 - val\_loss: 0.1013 - val\_acc: 0.9737

Epoch 5/10

60000/60000 [=====] - 5s 92us/step - loss: 0.1353 - acc: 0.9605 - val\_loss: 0.0981 - val\_acc: 0.9736

Epoch 6/10

60000/60000 [=====] - 5s 90us/step - loss: 0.1188 - acc: 0.9655 - val\_loss: 0.0997 - val\_acc: 0.9751

Epoch 7/10

60000/60000 [=====] - 5s 91us/step - loss: 0.1013 - acc: 0.9692 - val\_loss: 0.0942 - val\_acc: 0.9794

Epoch 8/10

60000/60000 [=====] - 5s 89us/step - loss: 0.0933 - acc: 0.9719 - val\_loss: 0.1042 - val\_acc: 0.9758

Epoch 9/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0830 - acc: 0.9746 - val\_loss: 0.0941 - val\_acc: 0.9799

Epoch 10/10

60000/60000 [=====] - 5s 88us/step - loss: 0.0747 - acc: 0.9760 - val\_loss: 0.1098 - val\_acc: 0.9767

```
In [39]: score = model4.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 51us/step

```
In [40]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [41]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9767

```
In [0]:
```

**changing dropout rate to 0.2**

```
In [42]: model4 = Sequential()

model4.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model4.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model4.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))

model4.add(BatchNormalization())
model4.add(Dropout(0.2))

model4.add(Dense(output_dim, activation='softmax'))

print(model4.summary())
```

Layer (type)	Output Shape	Param #
=====		
dense_14 (Dense)	(None, 364)	285740
dense_15 (Dense)	(None, 128)	46720
dense_16 (Dense)	(None, 64)	8256
batch_normalization_5 (Batch Normalization)	(None, 64)	256
dropout_5 (Dropout)	(None, 64)	0
dense_17 (Dense)	(None, 10)	650
=====		
Total params: 341,622		
Trainable params: 341,494		
Non-trainable params: 128		
None		

```
In [0]: model4.compile(optimizer='adam',
                      loss='categorical_crossentropy',
                      metrics=['accuracy'])
```



```
In [44]: # Training the model
history = model4.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 105us/step - loss: 0.2361 - acc: 0.9319 - val\_loss: 0.1084 - val\_acc: 0.9648

Epoch 2/10

60000/60000 [=====] - 5s 89us/step - loss: 0.0951 - acc: 0.9717 - val\_loss: 0.1040 - val\_acc: 0.9687

Epoch 3/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0647 - acc: 0.9810 - val\_loss: 0.0983 - val\_acc: 0.9693

Epoch 4/10

60000/60000 [=====] - 5s 88us/step - loss: 0.0511 - acc: 0.9843 - val\_loss: 0.0790 - val\_acc: 0.9742

Epoch 5/10

60000/60000 [=====] - 5s 88us/step - loss: 0.0386 - acc: 0.9879 - val\_loss: 0.0710 - val\_acc: 0.9785

Epoch 6/10

60000/60000 [=====] - 5s 90us/step - loss: 0.0332 - acc: 0.9896 - val\_loss: 0.0875 - val\_acc: 0.9754

Epoch 7/10

60000/60000 [=====] - 5s 87us/step - loss: 0.0248 - acc: 0.9918 - val\_loss: 0.0862 - val\_acc: 0.9765

Epoch 8/10

60000/60000 [=====] - 5s 88us/step - loss: 0.0235 - acc: 0.9921 - val\_loss: 0.0787 - val\_acc: 0.9793

Epoch 9/10

60000/60000 [=====] - 5s 87us/step - loss: 0.0196 - acc: 0.9939 - val\_loss: 0.0858 - val\_acc: 0.9794

Epoch 10/10

60000/60000 [=====] - 5s 88us/step - loss: 0.0195 - acc: 0.9937 - val\_loss: 0.0718 - val\_acc: 0.9813

```
In [45]: score = model4.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 55us/step

```
In [46]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# list of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [47]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9813

## Model with 5 hidden layers

```
In [0]: model5 = Sequential()

model5.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model5.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model5.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))

#model2.add(BatchNormalization())
#model2.add(Dropout(0.5))

model5.add(Dense(output_dim, activation='softmax'))

print(model5.summary())
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_36 (Dense)	(None, 364)	285740
dense_37 (Dense)	(None, 128)	46720
dense_38 (Dense)	(None, 96)	12384
dense_39 (Dense)	(None, 64)	6208
dense_40 (Dense)	(None, 32)	2080
dense_41 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 353,462		
Trainable params: 353,462		
Non-trainable params: 0		
None		

```
In [0]: model5.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
history = model5.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

Epoch 2/20

Epoch 3/20

Epoch 4/20

Epoch 5/20

Epoch 6/20

Epoch 7/20

Epoch 8/20

Epoch 9/20

Epoch 10/20

Epoch 11/20

Epoch 12/20

Epoch 13/20

Epoch 14/20

Epoch 15/20

Epoch 16/20

Epoch 17/20

Epoch 18/20

Epoch 19/20

Epoch 20/20

```
In [0]: score = model5.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 61us/step

```
In [0]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9804

## Model with 5 hidden layers along with BN and Dropout

```
In [0]: model6 = Sequential()

model6.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model6.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))

model6.add(BatchNormalization())
model6.add(Dropout(0.5))

model6.add(Dense(output_dim, activation='softmax'))

print(model6.summary())
```

Layer (type)	Output Shape	Param #
dense_48 (Dense)	(None, 364)	285740
dense_49 (Dense)	(None, 128)	46720
dense_50 (Dense)	(None, 96)	12384
dense_51 (Dense)	(None, 64)	6208
dense_52 (Dense)	(None, 32)	2080
batch_normalization_6 (Batch Normalization)	(None, 32)	128
dropout_6 (Dropout)	(None, 32)	0
dense_53 (Dense)	(None, 10)	330
Total params: 353,590		
Trainable params: 353,526		
Non-trainable params: 64		
None		

```
In [0]: model6.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [0]: # Training the model
history = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 30, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

Epoch 2/20

Epoch 3/20

Epoch 4/20

Epoch 5/20

Epoch 6/20

Epoch 7/20

Epoch 8/20

Epoch 9/20

Epoch 10/20

Epoch 11/20

Epoch 12/20

Epoch 13/20

Epoch 14/20

Epoch 15/20

Epoch 16/20

Epoch 17/20

Epoch 18/20

Epoch 19/20

Epoch 20/20

```
In [0]: score = model6.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 83us/step

```
In [0]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [0]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9814

## Changing Dropout to 0.8



```
In [48]: model6 = Sequential()

model6.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model6.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))

model6.add(BatchNormalization())
model6.add(Dropout(0.8))

model6.add(Dense(output_dim, activation='softmax'))

print(model6.summary())
```

Layer (type)	Output Shape	Param #
dense_18 (Dense)	(None, 364)	285740
dense_19 (Dense)	(None, 128)	46720
dense_20 (Dense)	(None, 96)	12384
dense_21 (Dense)	(None, 64)	6208
dense_22 (Dense)	(None, 32)	2080
batch_normalization_6 (Batch Normalization)	(None, 32)	128
dropout_6 (Dropout)	(None, 32)	0
dense_23 (Dense)	(None, 10)	330
Total params: 353,590		
Trainable params: 353,526		
Non-trainable params: 64		
None		

```
In [0]: model6.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [50]: # Training the model
history = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epoch,
                    verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10  
Epoch 2/10  
Epoch 3/10  
Epoch 4/10  
Epoch 5/10  
Epoch 6/10  
Epoch 7/10  
Epoch 8/10  
Epoch 9/10  
Epoch 10/10

```
In [51]: score = model6.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 56us/step

```
In [52]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [53]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9774

In [0]:

## Changing dropout to 0.2

```
In [54]: model6 = Sequential()

model6.add(Dense(364, activation='relu', input_shape=(input_dim,), kernel_initializer='random_uniform'))
model6.add(Dense(128, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(96, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(64, activation='relu', kernel_initializer='random_uniform'))
model6.add(Dense(32, activation='relu', kernel_initializer='random_uniform'))

model6.add(BatchNormalization())
model6.add(Dropout(0.2))

model6.add(Dense(output_dim, activation='softmax'))

print(model6.summary())
```

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_24 (Dense)	(None, 364)	285740
dense_25 (Dense)	(None, 128)	46720
dense_26 (Dense)	(None, 96)	12384
dense_27 (Dense)	(None, 64)	6208
dense_28 (Dense)	(None, 32)	2080
batch_normalization_7 (Batch Normalization)	(None, 32)	128
dropout_7 (Dropout)	(None, 32)	0
dense_29 (Dense)	(None, 10)	330
=====	=====	=====
Total params: 353,590		
Trainable params: 353,526		
Non-trainable params: 64		
None		

```
In [0]: model6.compile(optimizer='adam',
                        loss='categorical_crossentropy',
                        metrics=['accuracy'])
```

```
In [57]: # Training the model
history = model6.fit(X_train, y_train, batch_size = batch_size, epochs=nb_epochs, verbose= 1, validation_data=(X_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/10

60000/60000 [=====] - 6s 96us/step - loss: 0.1093 - acc: 0.9702 - val\_loss: 0.1149 - val\_acc: 0.9673

Epoch 2/10

60000/60000 [=====] - 6s 94us/step - loss: 0.0770 - acc: 0.9781 - val\_loss: 0.0869 - val\_acc: 0.9751

Epoch 3/10

60000/60000 [=====] - 6s 96us/step - loss: 0.0563 - acc: 0.9838 - val\_loss: 0.0875 - val\_acc: 0.9762

Epoch 4/10

60000/60000 [=====] - 6s 97us/step - loss: 0.0462 - acc: 0.9865 - val\_loss: 0.0839 - val\_acc: 0.9766

Epoch 5/10

60000/60000 [=====] - 6s 95us/step - loss: 0.0383 - acc: 0.9890 - val\_loss: 0.0804 - val\_acc: 0.9792

Epoch 6/10

60000/60000 [=====] - 6s 95us/step - loss: 0.0338 - acc: 0.9899 - val\_loss: 0.0968 - val\_acc: 0.9739

Epoch 7/10

60000/60000 [=====] - 6s 93us/step - loss: 0.0286 - acc: 0.9916 - val\_loss: 0.0852 - val\_acc: 0.9787

Epoch 8/10

60000/60000 [=====] - 6s 94us/step - loss: 0.0280 - acc: 0.9916 - val\_loss: 0.0806 - val\_acc: 0.9795

Epoch 9/10

60000/60000 [=====] - 6s 94us/step - loss: 0.0242 - acc: 0.9927 - val\_loss: 0.0972 - val\_acc: 0.9763

Epoch 10/10

60000/60000 [=====] - 6s 95us/step - loss: 0.0213 - acc: 0.9934 - val\_loss: 0.0747 - val\_acc: 0.9819

```
In [59]: score = model6.evaluate(X_test, y_test)
```


10000/10000 [=====] - 1s 65us/step

```
In [60]: fig,ax = plt.subplots(1,1)
ax.set_xlabel('Epochs') ; ax.set_ylabel('Loss')

# List of epoch numbers
list_of_epoch = list(range(1,nb_epoch+1))

train_loss = history.history['loss']
val_loss = history.history['val_loss']

ax.plot(list_of_epoch, val_loss, 'g', label="Validation Loss")
ax.plot(list_of_epoch, train_loss, 'r', label="Training Loss")
plt.legend()
plt.grid()
plt.show();
```



```
In [58]: print('Test accuracy:', score[1])
```

Test accuracy: 0.9774

## Conclusions

```

In [64]: # Please compare all your models using Prettytable Library
from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Numer of Layers", "BN", "Dropout", "Accuracy"]

x.add_row(["2", 'NO', "NO", 0.86])
x.add_row(["2", 'YES', 0.5, 0.98])
x.add_row(["2", 'YES', 0.8, 0.98])
x.add_row(["2", 'YES', 0.2, 0.98])

x.add_row(["3", 'NO', "NO", 0.87])
x.add_row(["3", 'YES', 0.5, 0.98])
x.add_row(["3", 'YES', 0.8, 0.97])
x.add_row(["3", 'YES', 0.2, 0.98])

x.add_row(["5", 'NO', 'NO', 0.98])
x.add_row(["5", 'YES', 0.5, 0.98])
x.add_row(["5", 'YES', 0.8, 0.97])
x.add_row(["5", 'YES', 0.2, 0.97])

print(x)

```

Numer of Layers	BN	Dropout	Accuracy
2	NO	NO	0.86
2	YES	0.5	0.98
2	YES	0.8	0.98
2	YES	0.2	0.98
3	NO	NO	0.87
3	YES	0.5	0.98
3	YES	0.8	0.97
3	YES	0.2	0.98
5	NO	NO	0.98
5	YES	0.5	0.98
5	YES	0.8	0.97
5	YES	0.2	0.97

In [0]: